

Adrian Erlikhman

(213) 330-6884 · erlikhman.adrian@gmail.com · adrianerlikhman.is-a.dev · [LinkedIn](#) · [GitHub](#)

PROFESSIONAL EXPERIENCE

Kiddom – Machine Learning Intern · K–12 EdTech 2026

- On the data science team with Director of Data Science Flora Xu, applying machine learning to educational content and product-usage data: turning classroom usage into models the product team can act on, and keeping them honest in production.

Fjord Venture Capital – Investment Intern · AI-focused venture fund 2024 – Present

- Ran due diligence on 20+ early-stage AI startups and wrote 10+ full investment memos (market sizing, competitive positioning, technical architecture, founder evaluation) for partner-level decisions.
- Led the screen of Y Combinator’s Summer 2025 batch with Managing Director Tatiana Janko; drafted the fund’s monthly LP updates.

Alliance for SoCal Innovation – Venture Outreach & Pipeline Intern 2024 – 2026

- Built the Alliance’s 250+ firm investor database in HubSpot and the schema it still sorts on, and led the backend redesign of the pipeline with the Chief Innovation Officer.
- Rewrote the founder intake form, routed companies to partners including Expert Dojo and the Long Beach Accelerator, and sat in on the SVP Life Science pilot committee’s pitch reviews.

VENTURES & CIVIC TECHNOLOGY

AIML-LI (AI/ML Literacy Initiative) – Co-Founder & Lead · aiml-initiative.org · grew out of the LACES AIML Club 2023 – Present

- Co-wrote a free AI-literacy course any teacher can run in twelve 30-minute sessions on Chromebooks: students train and break real models, study bias in real cases and pitch an AI project of their own. Taught in LACES AP Computer Science, spring 2026.
- Took it and a 12-week version to LAUSD’s Ed Tech & Innovation office (Sep 2026), which is sending them to the district’s Chief Academic and Chief Information Officers toward a standing AI-literacy elective. Advised by Philo Juang (Google DeepMind) and Amir Zohrenejad (Data Herald, YC W21); mentored by aiEDU.
- Organizing the LA Student AI Summit & Hackathon for 200+ LA public-school students: keynotes from Google DeepMind, Microsoft, Google and Snap, backed by Snap, Jane Street, MongoDB, Confluent, Discovery Education and Hacker Fund.

SafeJew – Co-Founder · safejew.org · JFEDLA Teen Innovation Grant 2024 – Present

- Shipped a live map of antisemitic incidents, in production at safejew.org: 800 across Greater LA, 100 of them individually sourced with a citable link on every pin, plus 594 on 33 university campuses.
- Wrote a dependency-free LA gazetteer that geocodes the places named in news prose, so a scheduled GitHub Action keeps the map current. Next.js, TypeScript and Supabase; rebuilt as V2 for a Hillel International pitch.

Safe Routes to School – Co-Creator · safe-routes-la.github.io 2026

- 3rd place internationally, Code for Transportation** (Young Coders’ Sphere). Scores 431,599 LA blocks for street-crime risk from 85,634 filtered LAPD incidents, then runs A* entirely in the browser, with no server.
- Found that the afternoon walk home is a student’s riskiest hour, and that riding the bus cuts exposure by up to 80% and is faster.

RESEARCH

Self-Attribution Is Not Self-Recognition: A Peer Baseline for LLM Judges

First author · with R. Erlikhman, M. Tarekegn · advised by Philo Juang (Google DeepMind) *Under review, JUDGE @ NeurIPS 2026*

- Introduced a peer baseline that scores each LLM judge’s self-attribution rate against what the other judges get on the same text, separating self-recognition from text that is simply easy to identify.
- Five frontier judges over 190 responses: Claude names its own writing 86.8% of the time against 53.9% for its peers (+32.9 points, Fisher $p = 1.6 \times 10^{-4}$), and an 18-feature stylometric classifier (86.3%) beats every judge.

Do LLM Fingerprints Survive Outside English?

First author · with M. Tarekegn, R. Erlikhman · advised by Philo Juang (Google DeepMind) *Poster submitted, IEEE MIT URTC 2026*

- 840 responses from five frontier models in seven languages, English to Hindi, attributed with 21 Universal Dependencies features: 59–72% accuracy in every language (chance is 20%), no drop for low-resource languages, and it survives machine translation.
- Across 4,195 judge calls, four of the five models attribute authorship at chance, while Claude’s self-recognition holds even in Hindi. [\[abstract\]](#) [\[poster\]](#) [\[code\]](#)

Proximity, Not Transmission: What the Patent Record Can and Cannot Show About NLP and Surveillance

First author · with Michael Tarekegn *Submitting, AI for Peace @ NeurIPS 2026*

- Linked 127,851 ACL Anthology papers to the 9,048 USPTO patents that cite them, and coded what each patent is for with a rubric that separates a system’s operator from its subject; revealing the applicant flips 10 of 31 non-neutral labels.
- Name-blind coding of 2,374 patents finds 0.46% surveillance, no standout subfield, no trend and no detectable difference from science-citing patents of the same era (Mantel–Haenszel OR 1.03, 95% CI 0.38–2.78). [\[abstract\]](#) [\[code\]](#)

How Robust Are the Legatum Prosperity Index Rankings?

A Monte Carlo and Sensitivity Analysis of Weighting, Normalization, and Aggregation

First author · with Ryan Erlikhman *Accepted with revisions, Journal of High School Science*

- Reproduced every published score for 167 countries exactly, then varied weights, normalization and aggregation: 10,000 draws under each of three weight priors and a 131,072-run Sobol’ design.

- Derived rank leverage, $c^2(1 - \rho^2)$, the only part of a weight's influence that can reorder countries; it tracks Sobol' total effects (Spearman 0.91–0.99). The ranking holds as a whole (median Spearman ≥ 0.9899) but not mid-table, where the median 90% rank interval spans 36 places. [\[pdf\]](#)

Does Portfolio Optimization Beat Splitting Evenly? Markowitz vs. the 1/N Portfolio

With Ryan Erlikhman

Under review, Oxford Journal of Student Scholarship

- Equal weights beat Markowitz in 74% of 250 calibrated simulations (Sharpe 0.70 vs. 0.51). The loss traces to estimating expected returns, which shrinkage fixes only with about two decades of stationary data. [\[pdf\]](#)

Advanced Mathematics Research · number theory, combinatorics, game theory; with PhD candidate Kirill Kovalenko 2022 – Present

COMPETITIONS & BUILDS

Ledger – 1st place, VISION HACK: South LA · Council District 9 Champion, U18 · Hacker Fund × City of LA Vision Lab Sep 2026

- Led the twelve-hour build across the vision pipeline, rule engine and frontend, and wrote the two-pass vision pipeline, the dashboard and renewal calendar, and the static fallback. Ledger puts the 18 programs a corner store answers to on one renewal calendar and scores a photo of a wholesale order against the SNAP stocking rule. [\[live\]](#) [\[code\]](#)

eDNAtlas – 1st place, Decode the Ocean Hackathon · Lovable × United Nations 2026

- Built and shipped the winning entry: a live map that turns environmental DNA into a plain-language health score for coastal sites, with the raw evidence and Darwin Core downloads. [\[live\]](#)

Citadel High School Terminal – Finalist (top 8) · Citadel | Citadel Securities, hosted by Correlation One 2026

- Wrote “dunerscore,” a Python tower-defense strategy that prices every wall, turret and attack lane each turn; iterated through replay post-mortems to the finals on a three-person team.

Y Combinator Startup School – Admitted · sub-5% acceptance, \$30,000+ in partner credits; validation to go-to-market 2026

More builds

- **Why Wrong**, misconception diagnosis for teachers (naive Bayes, then clustering; F1 0.87 on planted misconceptions) [\[live\]](#) · **Quant:** HMM regime-aware portfolio optimizer, LSTM forecaster, FinBERT sentiment signal, fraud detection (PR-AUC 0.36 vs. 0.28)

TALKS & PRESENTATIONS

NCSS 2026, National Consortium of STEM Schools – Invited Speaker, Student Panel on AI · San Diego Dec 2026

- Invited by aiEDU to its student panel in the pre-conference session “Transform Your School’s AI Readiness”: how students actually experience AI, and what adults need to know, for about 100 educators.

LAUSD Innovation Expo – Presenter · the AIML-LI curriculum, alongside Google, Microsoft and Canva

LAUSD Student Innovation Challenge – Invited Presenter · \$10,000 district challenge

EDUCATION

Los Angeles Center for Enriched Studies (LACES) – Class of 2027 · LAUSD magnet school Los Angeles, CA

- GPA 5.2 weighted / 4.0 unweighted · SAT 1520 (Math 770, EBRW 750) · 2× AP Scholar with Distinction
- 17 AP courses (11 completed, 6 this year)
- Dual enrollment: Linear Algebra (MATH 270, L.A. Southwest College) · Programming in Python (CS 119) · Python Programming for Data Science and Machine Learning (CS 121, planned)

LEADERSHIP & ATHLETICS

Competitive Fencing, Épée – USA Fencing A26, Team USA · LA International Fencing Center

- U.S. No. 49 Junior (U20) and No. 78 Senior; peak No. 17 Cadet (U17); No. 1 in Region 4. Team USA at the 2026 Cadet Pan American Championships in Bogotá (6th). In active recruitment with MIT.
- Top 32 Division I, National Championships · Top 8 Senior Team, Summer Nationals · Top 8 Cadet, SJCC · 4× Junior regional champion · 3× All-American and 3× All-Academic First Team

Jewish Student Union – President · 50+ members 2023 – Present

- Weekly programming with NCSY; organized a school-wide Holocaust Remembrance assembly featuring a survivor, with StandWithUs.

STEMsters – Co-Founder & President · 30+ volunteers, 4 elementary schools 2022 – Present

- Places LACES volunteers in four elementary schools to teach hands-on coding, robotics and engineering to underserved kids.

HONORS

- National Merit Semifinalist (about 16,000 of 1.3M+ entrants) · JFEDLA Teen Innovation Grant · SIFMA Stock Market Game, top 5% statewide as a solo entrant · MTAC piano, Level 9 candidate, five state-honors exams

SKILLS

- **Code:** Python, TypeScript/JavaScript, SQL, LaTeX; Next.js, React, Supabase/PostgreSQL, GitHub Actions, Git
- **ML & statistics:** PyTorch, scikit-learn, TensorFlow/Keras, pandas, NumPy, SHAP, Stanza; LLM evaluation, stylometry, Monte Carlo and Sobol' sensitivity analysis, hypothesis testing, time series, HMMs, A* search
- **Certifications:** Certiport IT Specialist in Software Development, Cybersecurity, Python, Data Analytics, AI and HTML/CSS; Deep Learning Specialization (Andrew Ng, Coursera); NVIDIA Deep Learning Fundamentals; Kaggle × Google Generative AI Intensive
- **Languages:** English, Russian (fluent), Spanish (intermediate)