Adrian Erlikhman

(213) 330-6884 | erlikhman.adrian@gmail.com | adrianerlikhman.is-a.dev | LinkedIn | GitHub

Professional Experience

Firstness 2026 – Present

Software Engineering & ML Intern | firstness.ai | reporting directly to the founder

- Refining how the platform connects into the research stack, and building demos that make it easier to understand.

Kiddom 2026 – Present

Machine Learning Internship & Mentorship | K–12 EdTech

- Designing ML and NLP approaches to teacher-facing problems in K–12 assessment with Kiddom's Director of Data Science; built **Why Wrong**, which names the misconception behind each wrong answer on a class's answer sheets (F1 0.87).

Fjor Venture Capital 2024 – 2026

Investment & Software Engineering Intern | $15M early-stage fund

- Built scoring models to screen 134 startups in Y Combinator's Summer 2025 batch and led the screen with Managing Director Tatiana Janko; ran diligence on 20+ AI startups and wrote 10+ investment memos.

Alliance for SoCal Innovation 2024 – 2026

Venture & Data Strategy Intern

- Built a Python/SQL investor-matching engine over 10K contacts at 250+ firms; deduplication cut lookups from days to under a minute. Worked with SoCal deep-tech startups, including Caltech spinouts.

Ventures

AIML-LI (AI/ML Literacy Initiative) 2023 – Present

Co-Founder & Lead | aiml-initiative.org

- Built a free AI-literacy course any teacher can run cold on the Chromebooks a class already has: twelve 30-minute sessions plus a 12-week, six-unit version with Colab notebooks, instructor guides and assessments. Taught in LACES AP Computer Science, spring 2026.
- Took it to LAUSD's Ed Tech & Innovation office, which is carrying it to the district's Chief Academic and Chief Information Officers toward a standing AI-literacy elective.
- Organizing the LA Student AI Summit & Hackathon for 500+ public-school students: keynotes from Google DeepMind, Microsoft, Google and Snap, backed by Snap, Jane Street, MongoDB, Confluent and Discovery Education.
- **Invited by aiEDU:** speaker on its student panel on AI at NCSS 2026 (National Consortium of STEM Schools, San Diego, Dec 2026) and a guest on its podcast with CEO Alex Kotran.

SafeJew 2024 – Present

Co-Founder | safejew.org | funded by a JFEDLA Teen Innovation Grant

- Shipped a live map of 1,394 antisemitic incidents, 800 across Greater LA and 594 on 33 university campuses, each with its source; funded by a Teen Innovation Grant from the Jewish Federation of Greater Los Angeles and distributed through its Community Security Initiative.
- Wrote a dependency-free LA gazetteer that geocodes places named in news prose, so a daily GitHub Action keeps the map current with no human in the loop; rebuilt as V2 for a Hillel International pitch.

Research

Whose Line Is It Anyway? Auditing Self- and Other-Recognition in LLM Judges

Preprint, arXiv forthcoming

- *Co-first author, with Michael Tarekegn and Philo Juang.* Nine LLM judges, 1,270 texts: the usual self-recognition test flags 15 of 37 cases, 3 beat a peer baseline and none beats its strongest single peer; an 18-feature stylometric classifier (86.3%) beats every judge.

Interpretable Multilingual Attribution of Frontier LLMs

Best Paper Finalist, IEEE BigData 2026 HS Symposium

- *Co-first author.* 840 essays from five models in seven languages: attribution holds at 59–72% against 20% chance, with no detected drop for low-resource languages, and survives Google Translate.

Curated to Win: Certifying Reward-Model and Judge Selection Against Evaluation-Set Curation

Accepted, ARRSI Workshop @ AACL-IJCNLP 2026

- *Sole author.* A certificate for how much of a validation set can be removed before the chosen reward model or LLM judge changes: on RewardBench, 12 of 2,985 items.

How Robust Are the Legatum Prosperity Index Rankings?

In revision, Journal of High School Science

- *Co-first author, with Ryan Erlikhman.* Derived rank leverage, $c^2(1-\rho^2)$, which tracks Sobol' total effects (Spearman 0.91–0.99); the ranking holds globally but not mid-table, where the median 90% rank interval spans 36 places.

Does Portfolio Optimization Beat Splitting Evenly?

Published, Oxford Journal of Student Scholarship, Oct 2026

- *Co-first author, with Ryan Erlikhman.* Equal weights beat Markowitz in 74% of 250 calibrated simulations (Sharpe 0.70 against 0.51); the loss traces to estimating expected returns.

Competitions

- **Finalist, Citadel High School Terminal** (top 8 of 45): "dunderscore," a Python bot, revised after each round's replay review. about
- **Y Combinator Startup School** (sub-5%, $30,000+ in credits), admitted as AIML-LI's founder. about
- **1st place, VISION HACK: South LA** (Council District 9 Champion, U18). Ledger: led the build across the vision pipeline, rule engine and frontend, and wrote the two-pass vision pipeline, the dashboard and renewal calendar, and the static fallback. live
- **1st place, Finance track, ACM TrojanHacks 2026**. Monte: picks the best credit card at each purchase. live devpost
- **Winner, Agent Behaviour Monitoring track, SEPHacks 2026**. Hivemind: shared memory for AI agents. live devpost
- **3rd place internationally, Code for Transportation**. Safe Routes: A* over 431,599 crime-scored LA blocks. live
- **1st place, Decode the Ocean Hackathon** (Lovable × UN). eDNAtlas: live coastal health scores from environmental DNA. live

Education

Los Angeles Center for Enriched Studies (LACES) Los Angeles, CA

Class of 2027 | LAUSD magnet school

- GPA 5.2 weighted, 4.0 unweighted | SAT 1520 | 2× AP Scholar with Distinction | 17 APs
- Dual enrollment: Linear Algebra (MATH 270) | Programming in Python (CS 119) | Python for DS and ML (CS 121, planned)

Leadership & Athletics

- **Épée fencing:** USA Fencing A26; Team USA, 2026 Cadet Pan American Championships (6th); U.S. No. 49 Junior and No. 78 Senior, peak No. 17 Cadet; 3× All-American First Team; in active recruitment with MIT.
- **President**, Jewish Student Union (50+ members) | **Co-Founder & President**, STEMsters (30+ volunteers in 4 elementary schools)

Skills

- Python, TypeScript/JavaScript, SQL, LaTeX | PyTorch, scikit-learn, TensorFlow, pandas, NumPy | Next.js, React, Supabase | LLM evaluation, stylometry, Monte Carlo and Sobol' sensitivity analysis | English, Russian (fluent), Spanish